

Chapter 3: Processes


Process Concept

- An operating system executes a variety of programs:
 - Batch system – jobs
 - Time-shared systems – user programs or tasks
- Terms *job* and *process* almost interchangeably
- Process – a program in execution; process execution must progress in sequential fashion
 - It is a unit of work in a modern time-sharing system.
- A process includes:
 - program counter
 - Stack
 - Data section
 - Heap

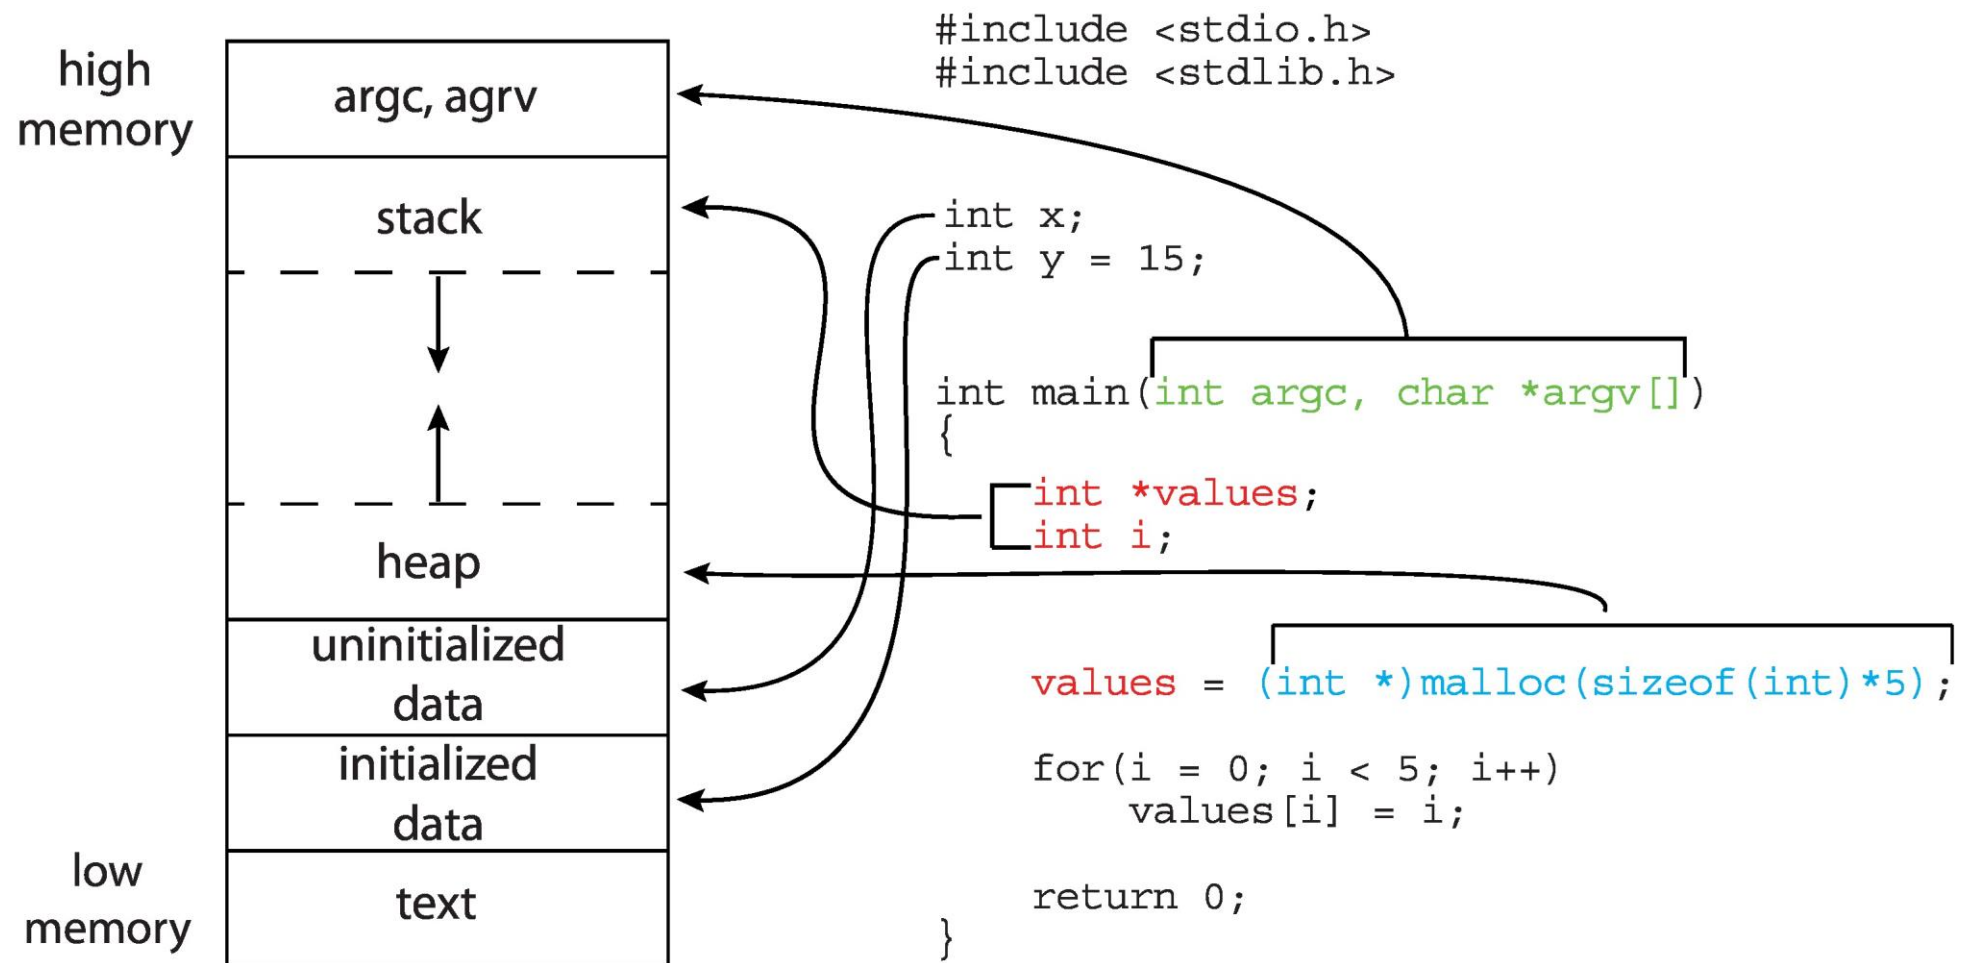
The Process

- Multiple parts
 - The program code, also called **text section**
 - Current activity including **program counter**, processor registers
 - **Stack** containing temporary data
 - ▶ Function parameters, return addresses, local variables
 - **Data section** containing global variables
 - **Heap** containing memory dynamically allocated during run time
- Program is passive entity, process is active
 - Program becomes process when executable file loaded into memory
- Execution of program started via GUI mouse clicks, command line entry of its name, etc
- One program can be several processes
 - Consider multiple users executing the same program

Process in Memory

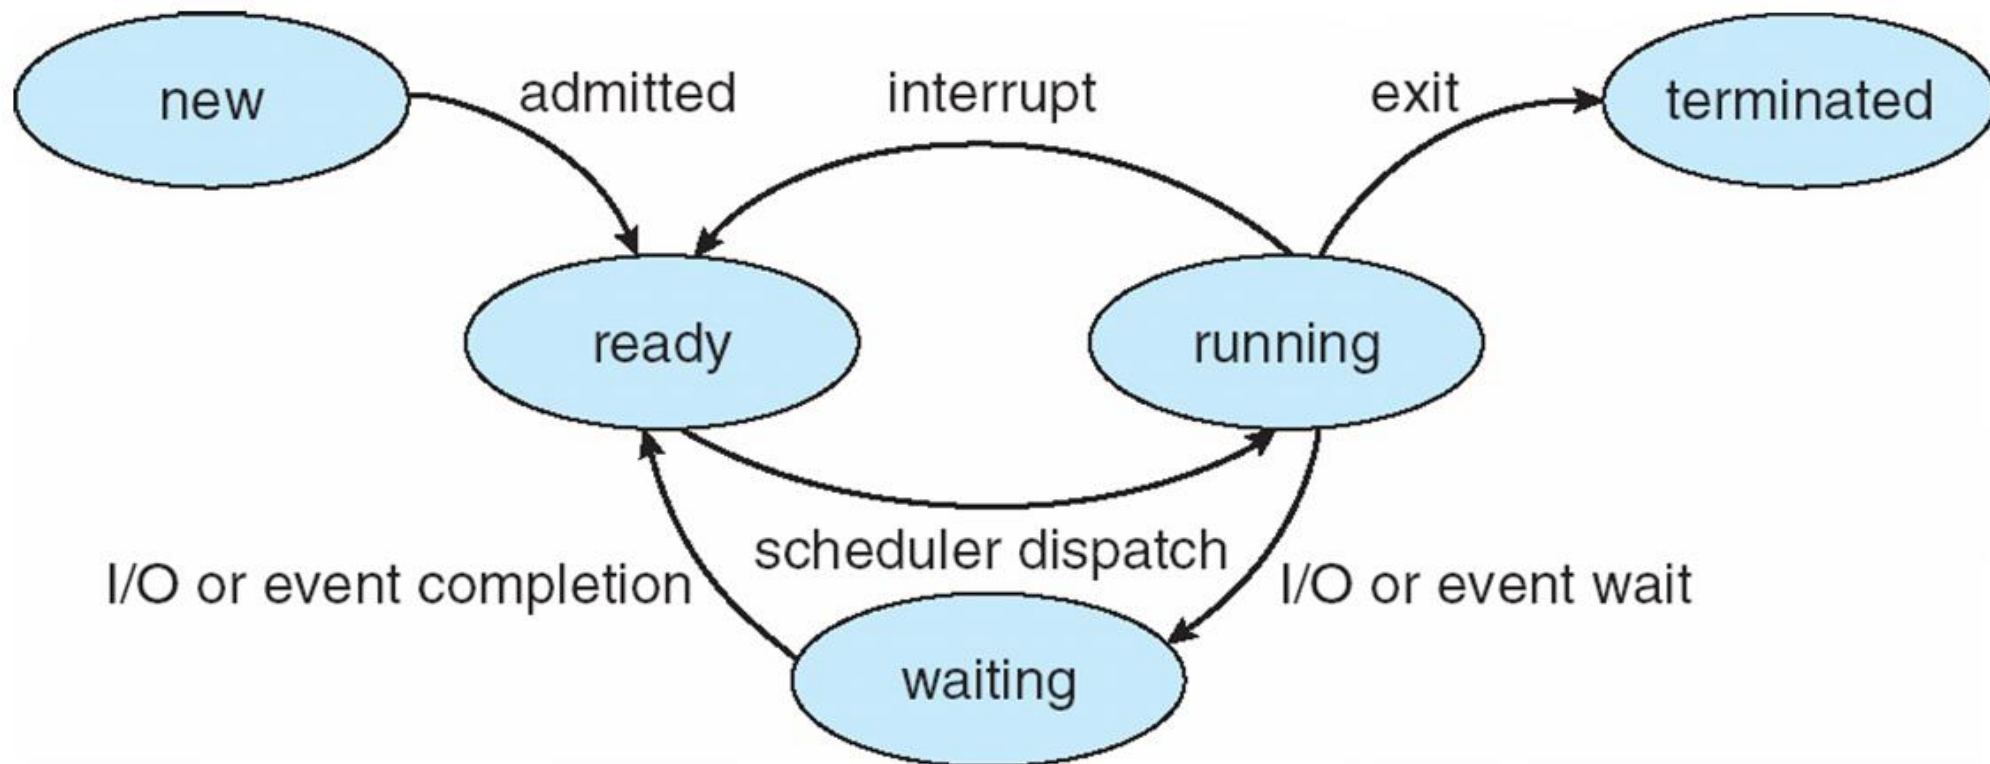


Memory Layout of a C Program



Process State

- As a process executes, it changes *state*
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a processor
 - **terminated**: The process has finished execution



Only one process can be running on any processor at any instant

Process Control Block (PCB)

Also called as a Task Control Block

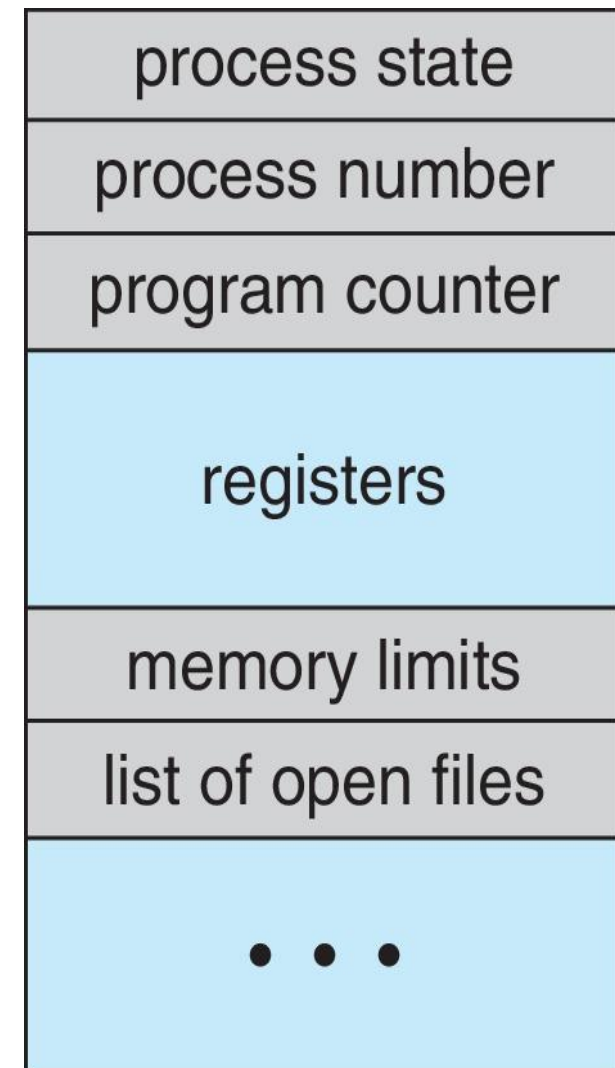
Information associated with each process

- ❑ Process state- **ready, running, etc.**
- ❑ Program counter- **keeps track of the next instruction to be executed**
- ❑ CPU registers- **accumulators, index registers, general purpose registers**
- ❑ CPU scheduling information- **process priority and scheduling queues**
- ❑ Memory-management information- **values of base-index registers, page table, segment tables**
- ❑ Accounting information- **amount of CPU time used, time limits, etc.**
- ❑ I/O status information- **list of I/O devices allocated to the process etc.**

Process Control Block (PCB)

Information associated with each process(also called **task control block**)

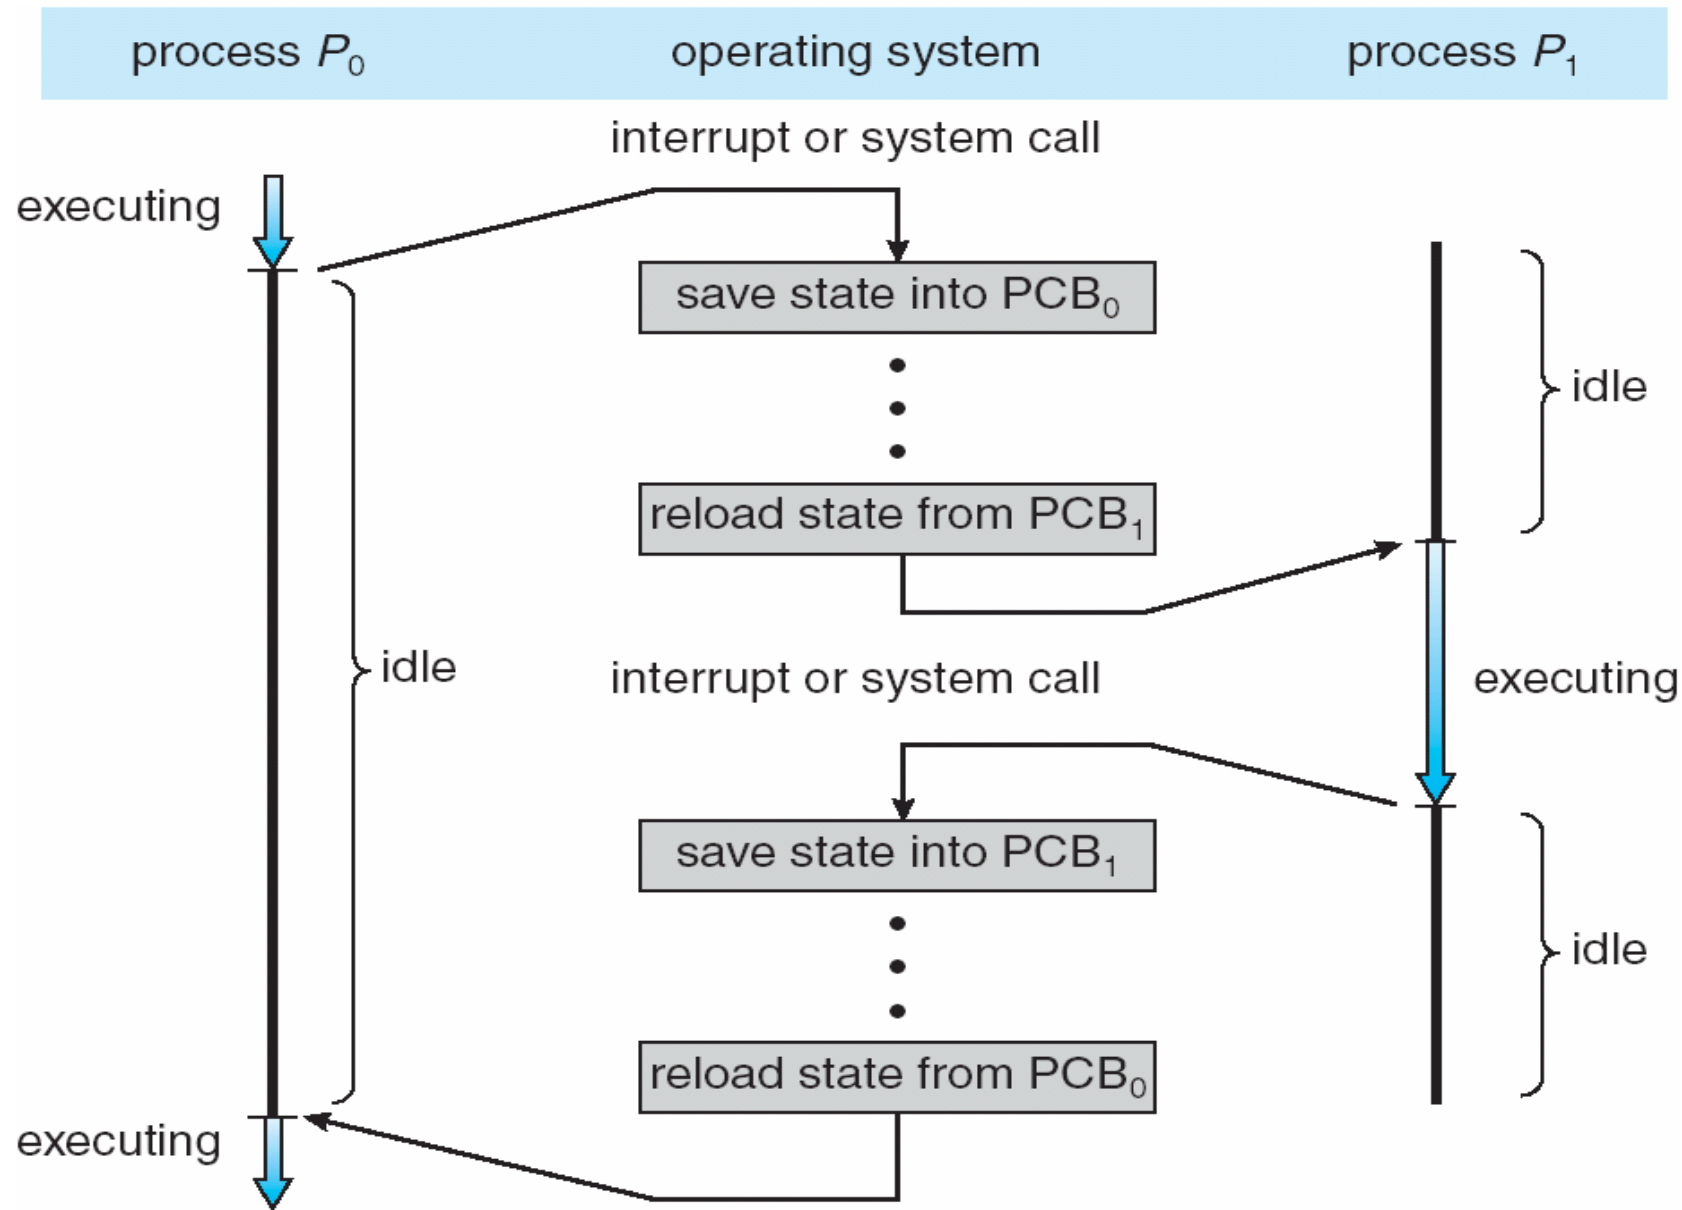
- ❑ Process state – running, waiting, etc.
- ❑ Program counter – location of instruction to next execute
- ❑ CPU registers – contents of all process-centric registers
- ❑ CPU scheduling information- priorities, scheduling queue pointers
- ❑ Memory-management information – memory allocated to the process
- ❑ Accounting information – CPU used, clock time elapsed since start, time limits
- ❑ I/O status information – I/O devices allocated to process, list of open files



Context Switch

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process via a **context switch**.
- **Context** of a process represented in the PCB
- Context-switch time is overhead; the system does no useful work while switching
 - The more complex the OS and the PCB -> longer the context switch

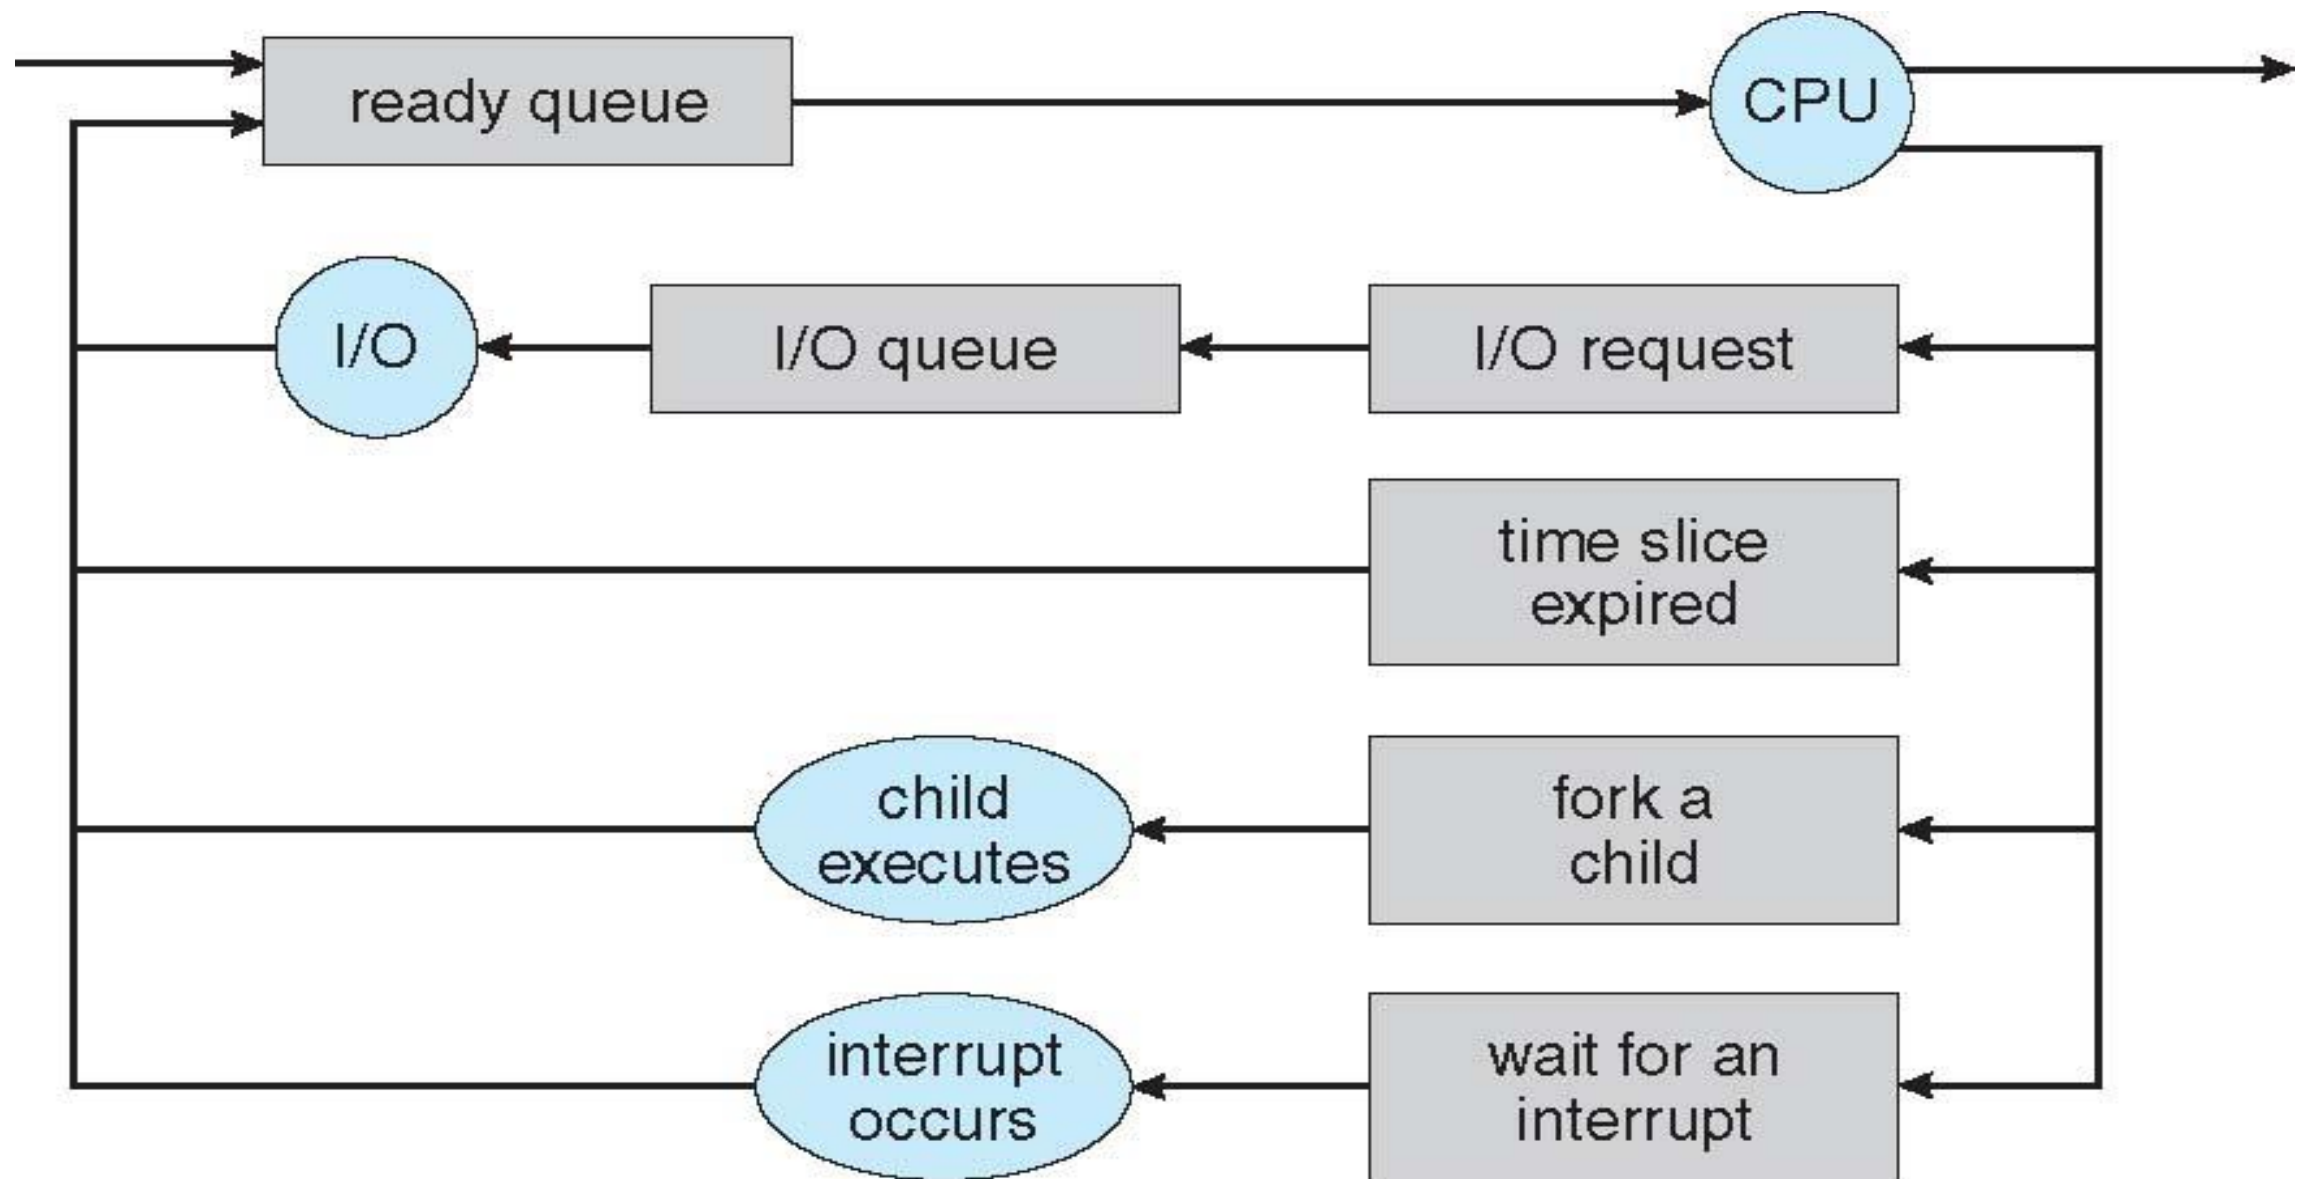
CPU Switch From Process to Process



Process Scheduling

- Maximize CPU use, quickly switch processes onto CPU for time sharing
- **Process scheduler** selects among available processes for next execution on CPU
- Maintains **scheduling queues** of processes
 - **Job queue** – set of all processes in the system
 - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
 - **Device queues** – set of processes waiting for an I/O device
 - ▶ Each device has its own device queue
 - Processes migrate among the various queues

Representation of Process Scheduling



Queueing Diagram

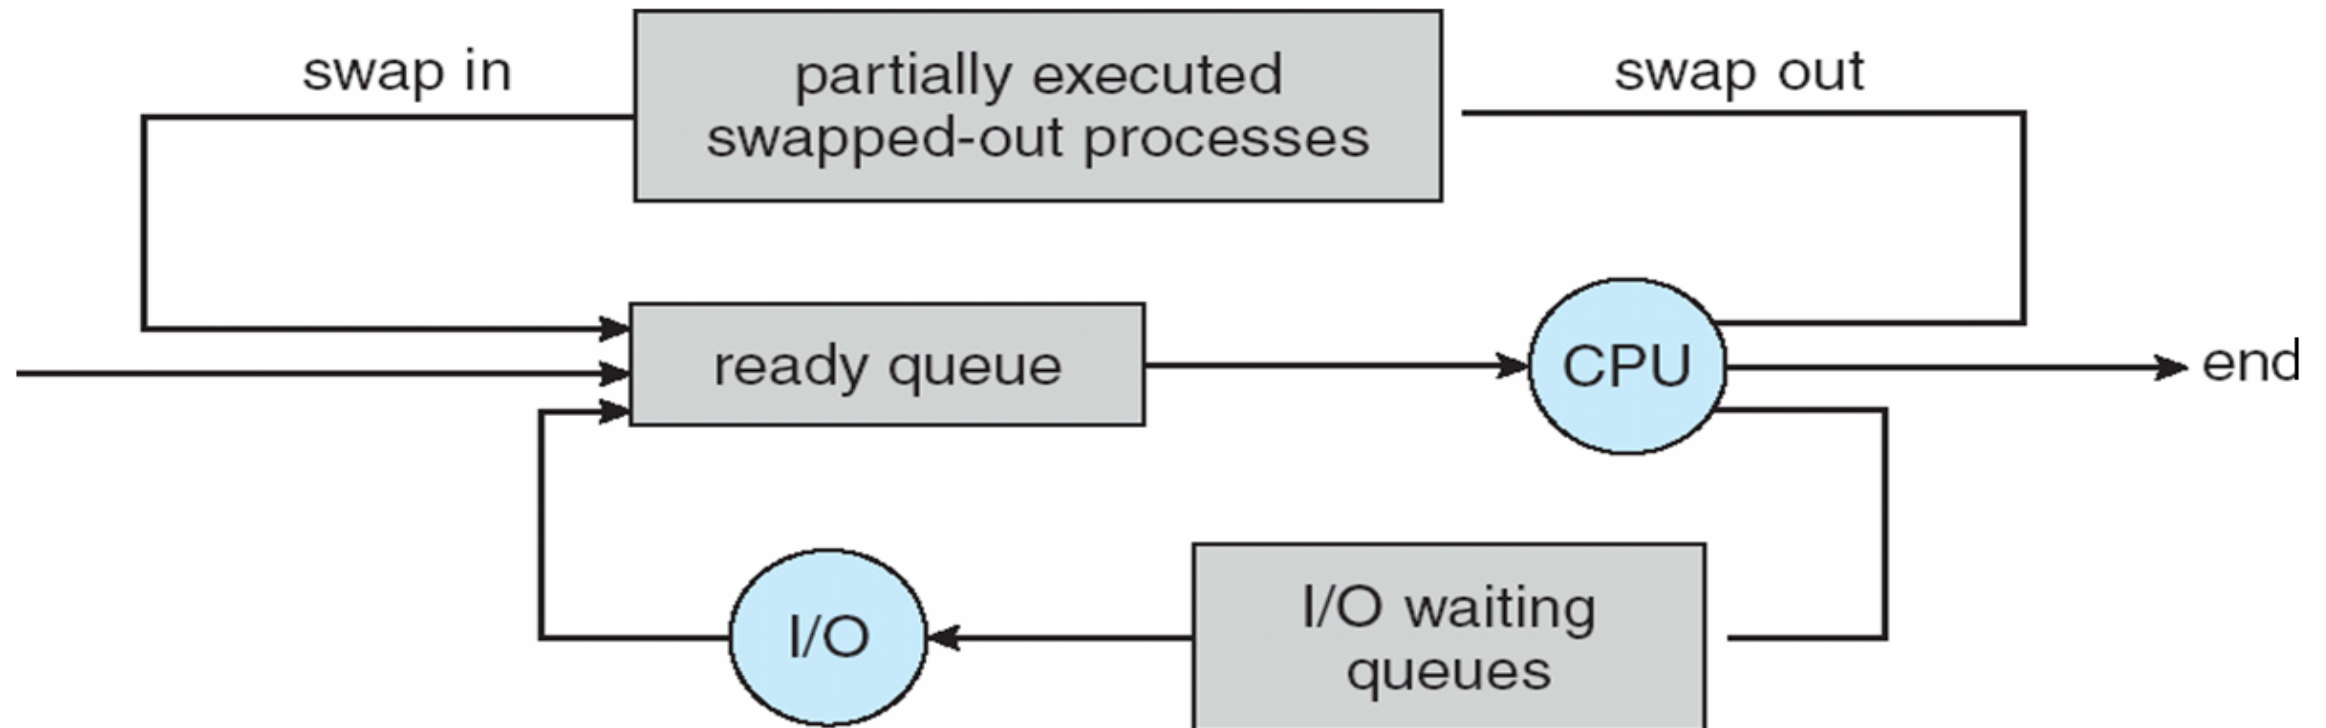
Schedulers

- **Long-term scheduler** (or job scheduler) – selects which processes should be brought into the ready queue
- **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system

Schedulers (Cont.)

- Short-term scheduler is invoked very frequently (milliseconds) \Rightarrow (must be fast)
- Long-term scheduler is invoked very infrequently (seconds, minutes) \Rightarrow (may be slow)
- The long-term scheduler controls the *degree of multiprogramming*
- Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; ~~few~~ **very long,** might have a few long CPU bursts

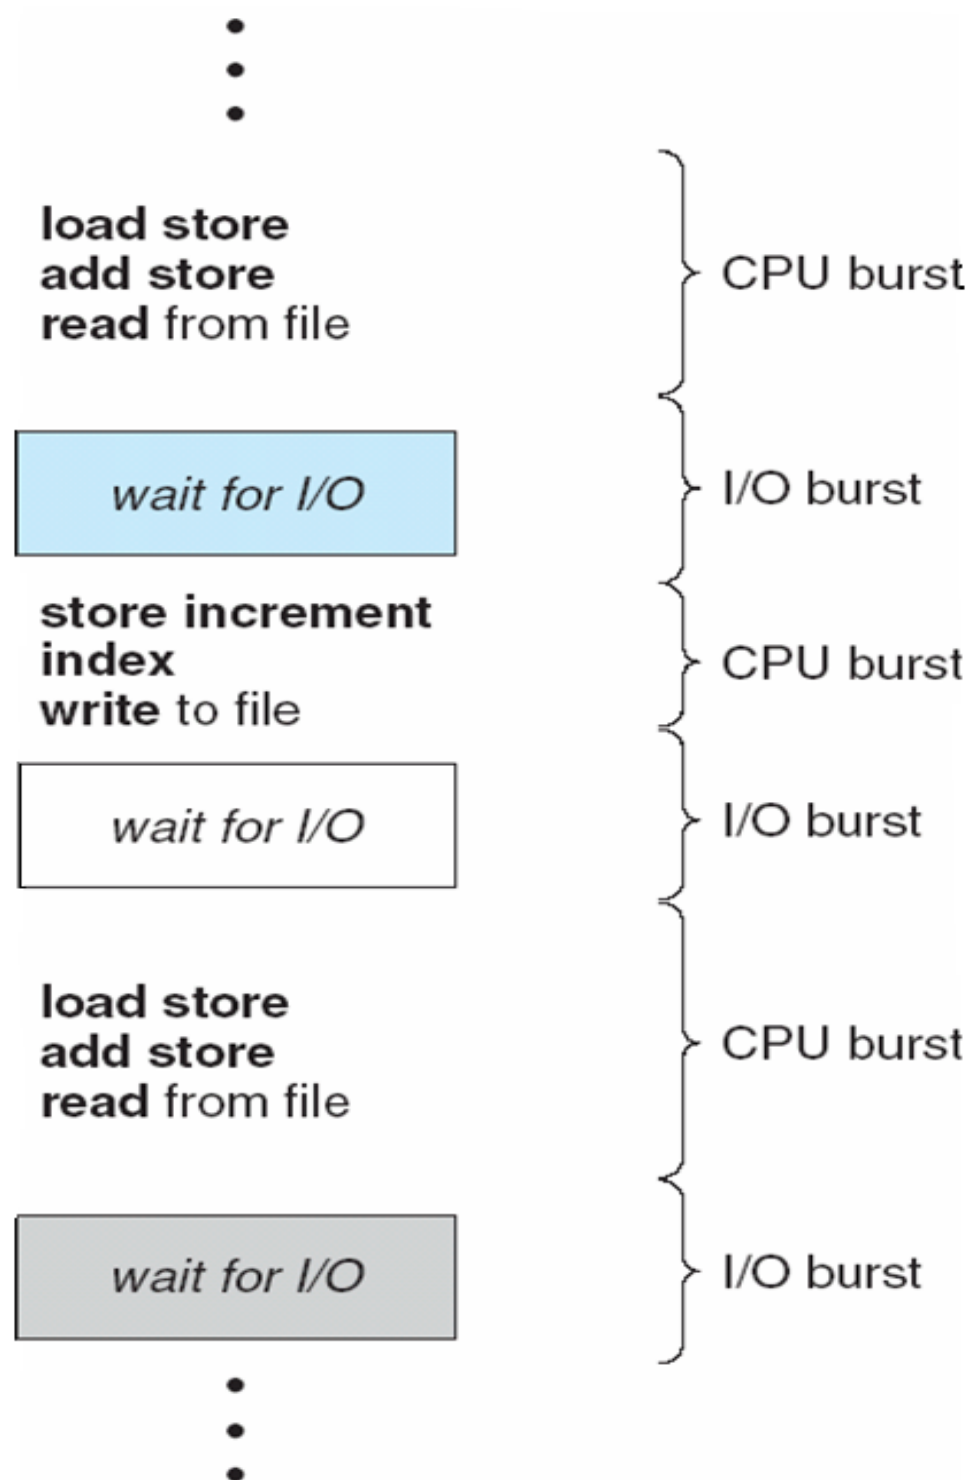
Addition of Medium Term Scheduling



- Key Idea


- Sometimes it is advantageous, to remove process from memory (active contention of CPU) to reduce the degree of multiprogramming.
- The process can be reintroduced into the ready queue and its execution can be continued where it left off. The scheme is called **swapping**.
- Swapping may be necessary to improve the process mix.

Alternating Sequence of CPU and I/O Bursts



- Multiprogramming
 - Several processes are kept in memory at one time.
 - When one process has to wait for I/O.
 - The OS takes CPU away from that process and gives the CPU to another process.

Histogram of CPU-burst Times



- Multiprogramming
 - Large number of short CPU bursts
 - Small number of Long CPU bursts
 - An I/O bound process has many short CPU bursts

CPU Scheduler

- ❑ Selects from among the processes in ready queue, and allocates the CPU to one of them
 - ❑ Queue may be ordered in various ways (Not in FIFO only)
- ❑ CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- ❑ Scheduling under 1 and 4 is nonpreemptive
- ❑ All other scheduling is preemptive
- ❑ No-preemptive- Once the CPU has been allocated to a process, the process keeps the CPU until termination or switching to waiting state.


Scheduling Criteria

- CPU utilization – keep the CPU as busy as possible
- Throughput – #number of processes that complete their execution per time unit
- Turnaround time – amount of time to execute a particular process (Time from submission to completion)
- Waiting time – amount of time a process has been waiting in the ready queue
- Response time – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Suppose that the processes arrive in the order: P_1 , P_2 , P_3
The Gantt Chart for the schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
□ Average waiting time: $(0 + 24 + 27)/3 = 17$

FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- Convoy effect - short process behind long process
 - Consider one CPU-bound and many I/O-bound processes

FCFS Scheduling (Cont.)

- FCFS scheduling algorithm is non-preemptive
 - Once the CPU has been allocated to a process, it keeps the CPU until process terminates or by requesting I/O.
 - FCFS is therefore troublesome for time sharing systems.
 - ▶ It would be disastrous to allow one process to keep the CPU for an extended period.


Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst
 - Use these lengths to schedule the process with the shortest time
 - if CPU bursts of two processes are same, then FCFS scheduling is used for selection
- SJF is optimal – gives minimum average waiting time for a given set of processes
 - The difficulty is knowing the length of the next CPU request
 - Could ask the user

Example of SJF

<u>Process</u>	<u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

□ SJF scheduling chart



□ Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$

Knowing the length of the next CPU request?

Determining Length of Next CPU Burst

- Can only estimate the length – should be similar to the previous one
 - Then pick process with shortest predicted next CPU burst
- Can be done by using the length of previous CPU bursts, using exponential averaging

1. t_n = actual length of n^{th} CPU burst
2. τ_{n+1} = predicted value for the next CPU burst
3. $\alpha, 0 \leq \alpha \leq 1$
4. Define :

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n.$$

- Commonly, α set to $1/2$
- Preemptive version called shortest-remaining-time-first

Examples of Exponential Averaging

□ $\alpha = 0$

□ $\tau_{n+1} = \tau_n$


□ Recent history does not count

□ $\alpha = 1$

□ $\tau_{n+1} = \alpha t_n$

□ Only the actual last CPU burst counts

Prediction of the Length of the Next CPU Burst




- Preemptive version called shortest-remaining-time-first

Example of Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

- *Preemptive* SJF Gantt Chart



- Average waiting time = $[(10-1)+(1-1)+(17-2)+5-3]/4 = 26/4 = 6.5$ msec


Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - Preemptive
 - Nonpreemptive
- SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- Problem \equiv Starvation – low priority processes may never execute
- Solution \equiv Aging – as time progresses increase the priority of the process

Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

□ Priority scheduling Gantt Chart



□ Average waiting time = 8.2 msec


Round Robin (RR)

- Each process gets a small unit of CPU time (time quantum q), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- Timer interrupts every quantum to schedule next process
- Performance
 - q large \Rightarrow FIFO
 - q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high
 - if time quantum is extremely small,
 - ▶ processor sharing

Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3


□ The Gantt chart is:



- Typically, higher average turnaround than SJF, but better *response*
- q should be large compared to context switch time
- q usually 10ms to 100ms, context switch < 10 usec

Time Quantum and Context Switch Time

Typically, higher average turnaround than SJF,
but better *response*

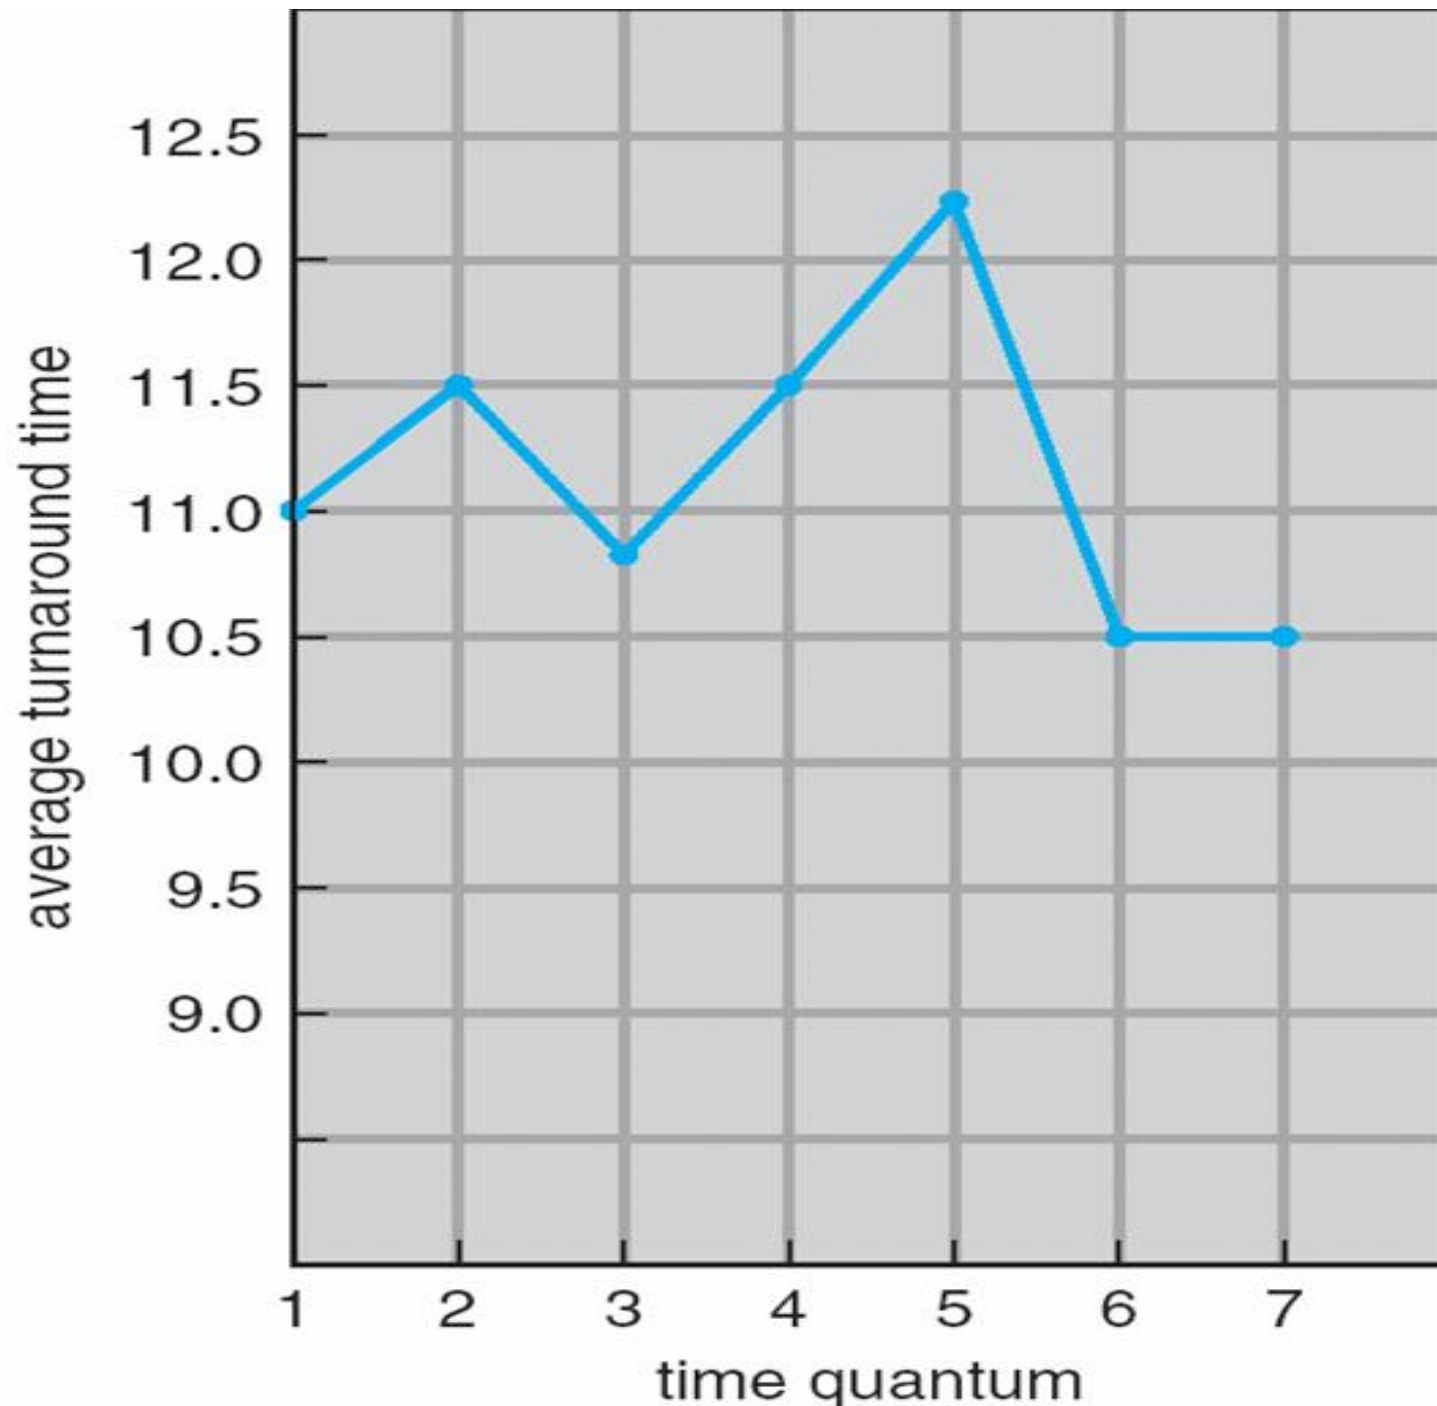


Turnaround Time Varies With The Time Quantum

process	time
P_1	6
P_2	3
P_3	1
P_4	7

Turnaround time with quantum =5

Turnaround Time Varies With The Time Quantum




process	time
P_1	6
P_2	3
P_3	1
P_4	7

80% of CPU
bursts should be
shorter than q

Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

□ The Gantt chart is:



- Typically, higher average turnaround than SJF, but better *response*
- q should be large compared to context switch time
- q usually 10ms to 100ms, context switch < 10 usec


<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	8
P_2	0.4	4
P_3	1.0	1


<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	2	2
P_2	1	1
P_3	8	4
P_4	4	2
P_5	5	3

<u>Process</u>	<u>Priority</u>	<u>Burst</u>	<u>Arrival</u>
P_1	40	20	0
P_2	30	25	25
P_3	30	25	30
P_4	35	15	60
P_5	5	10	100
P_6	10	10	105

The length of a time quantum is 10 units

If a process is preempted by a higher-priority process, the preempted process is placed at the end of the queue.





P1: $20 - 0 = 20$, P2: $80 - 25 = 55$, P3: $90 - 30 = 60$, P4: $75 - 60 = 15$, P5:
 $120 - 100 = 20$, P6: $115 - 105 = 10$

P1: 0, P2: 40, P3: 35, P4: 0, P5: 10, P6: 0

Algorithm Evaluation


- How to select CPU-scheduling algorithm for an OS?
- Determine criteria, then evaluate algorithms
- Deterministic modeling
 - Type of analytic evaluation
 - Takes a particular predetermined workload and defines the performance of each algorithm for that workload
- Consider 5 processes arriving at time 0:

<u>Process</u>	<u>Burst Time</u>
P_1	10
P_2	29
P_3	3
P_4	7
P_5	12


Deterministic Evaluation

- For each algorithm, calculate minimum average waiting time
- Simple and fast, but requires exact numbers for input, applies only to those inputs


- FCS is 28ms:



- Non-preemptive SFJ is 13ms:



- RR is 23ms:

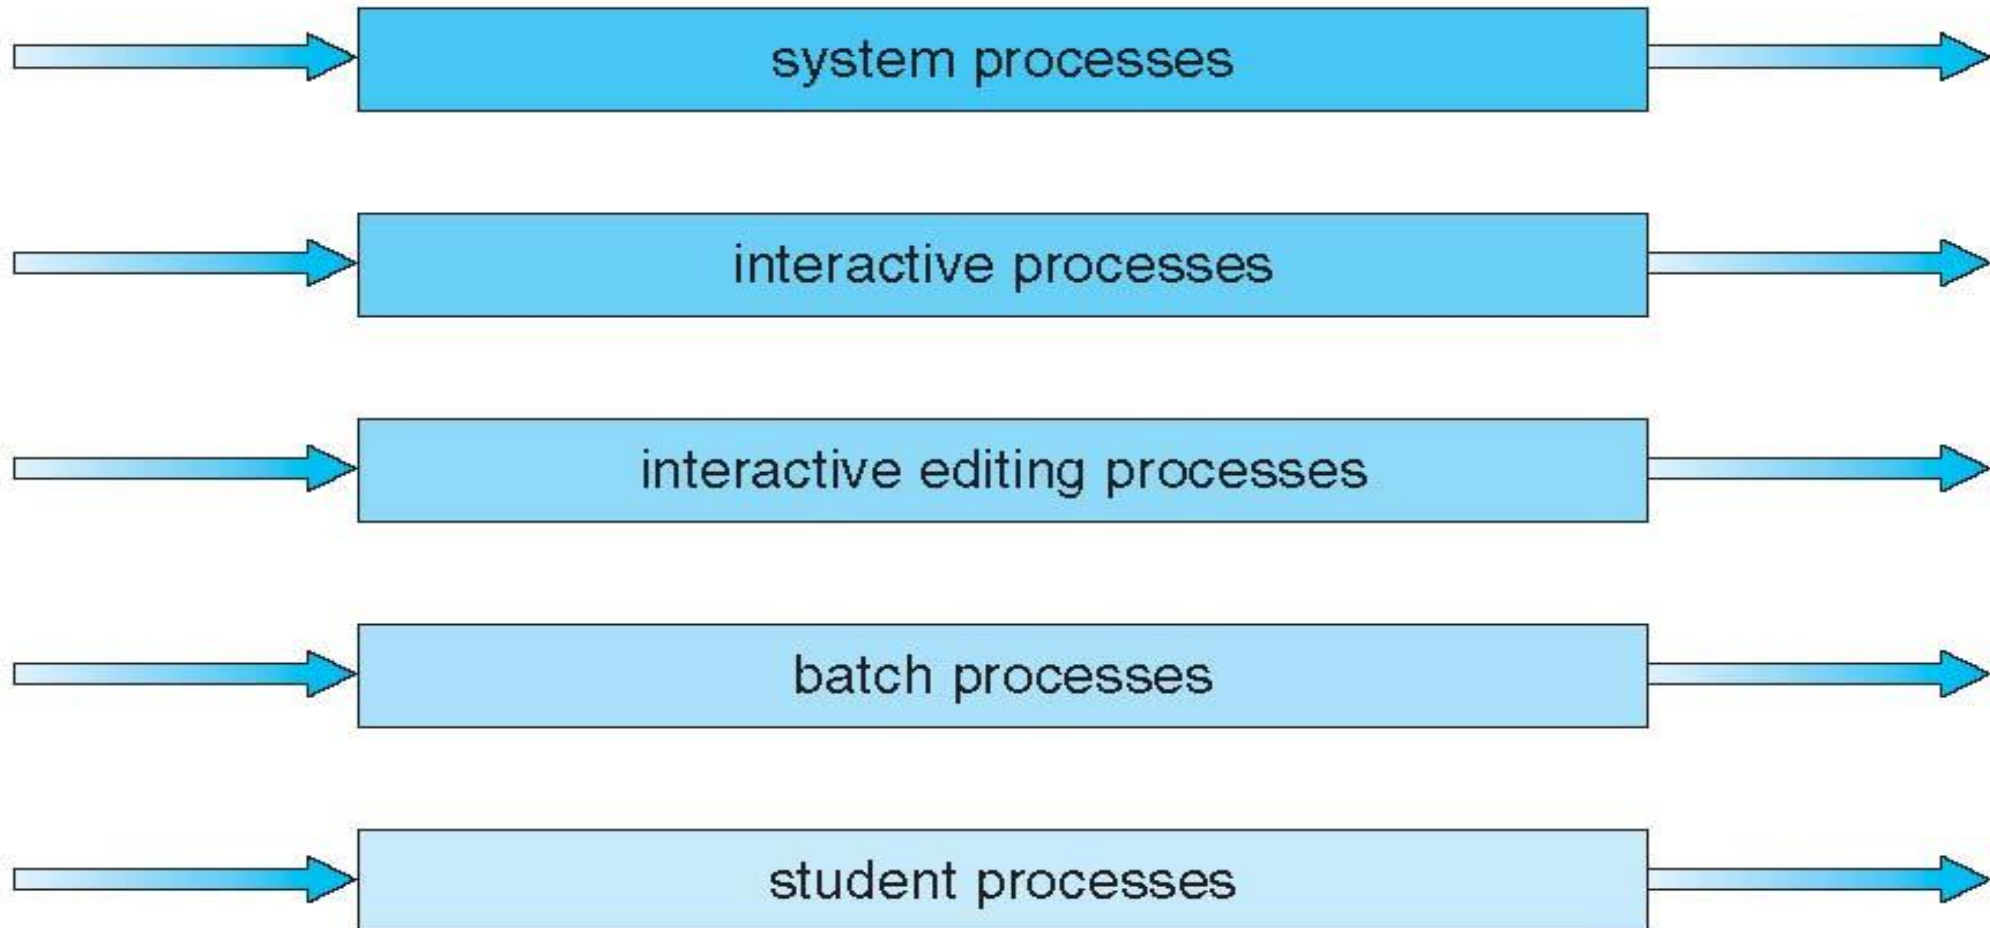


Multilevel Queue

- Ready queue is partitioned into separate queues, eg:
 - foreground (interactive)
 - background (batch)
- Process permanently in a given queue based on the property of the process
- Each queue has its own scheduling algorithm:
 - foreground – RR
 - background – FCFS
- Scheduling must be done between the queues:
 - Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
 - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR 20% to background in FCFS

Multilevel Queue Scheduling

highest priority



lowest priority

Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way
 - ▶ If process uses too much CPU time.
 - Moved to lower- priority queue
 - ▶ I/O bound and interactive process
 - Moved to higher priority queue
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service

Example of Multilevel Feedback Queue

- Three queues:
 - Q_0 – RR with time quantum 8 milliseconds
 - Q_1 – RR time quantum 16 milliseconds
 - Q_2 – FCFS

- Scheduling
 - A new job enters queue Q_0 which is served FCFS
 - ▶ When it gains CPU, job receives 8 milliseconds
 - ▶ If it does not finish in 8 milliseconds, job is moved to queue Q_1
 - At Q_1 job is again served FCFS and receives 16 additional milliseconds
 - ▶ If it still does not complete, it is preempted and moved to queue Q_2

Multilevel Feedback Queues

